

Supplementary Material for “Silent Dissent: LLM Agents That Yield to the Majority Still Represent Their Original Premise”

Ziang Ni*

Delft University of Technology
Delft, The Netherlands
ziangni@tudelft.nl

Peng Zou*

Sun Yat-sen University
Guangzhou, China
zoup7@mail2.sysu.edu.cn

Abstract

This supplement documents the pre-registration record, the prompts, the data and readout details, the full result tables behind the paper, the own-answer addendum, the exploratory robustness analyses and the negative results, the compute used, and the use of AI tools. Everything here is additional detail; the paper is self-contained. Section references of the form “paper §*n*” point to the main paper.

Keywords

Multi-agent debate; Conformity; LLM agents; Pre-registration; Jacobian lens

1 Contents of the Repository

The repository <https://github.com/ziangni-sys/silent-dissent> contains this document, the paper and the code:

- `silent-dissent/`: the library (prompts and conditions, lenses and the J-lens estimator, readouts, injection, metrics).
- `scripts/`: one command-line script per step (fact check and split, development runs, selection and registration, test runs, timeline, injection, free debate and its follow-ups, analysis, cross-model summary, paper figures and tables).
- `configs/`: the configuration of every model and lens refit.
- `prereg/`: all pre-registration files, byte-identical to the committed versions (Section 2).
- `tests/`: unit and end-to-end tests that run the whole pipeline on a small random model without downloads.
- `docs/results/crossmodel_summary.json`: the summary from which every number, table and figure of the paper is computed; `scripts/entity_paper.py` and `scripts/entity_supp_tables.py` regenerate the figures and tables from it on a CPU.
- The exploratory robustness analyses (Section 8), computed from the raw test readouts by `scripts/entity_robust.py`: `docs/results/robust_summary.json`.
- The own-answer addendum’s verdicts and summaries per model (Section 7): `docs/results/own/`.

Raw per-record outputs (lens ranks for every record, position, layer and role; generated answers; logs) are much larger and are not included; they will be released with the camera-ready version.

2 Pre-registration Record

Procedure. For every model the pipeline was: (1) check which facts the model knows and split them by bridge entity into development (30% of bridges) and test sets; (2) run the validation conditions on the development split only; (3) run the selection script, which applies the declared rules to the development data and writes a registration file (format, layer band, categories, hypotheses, validation results and every development-stage adjustment); (4) commit and push the file; (5) only then run the test split. The selection script refuses to overwrite an existing registration. Follow-up questions were registered as addenda under the same procedure, with their rules declared in the configuration before the development runs.

Rules. Both band rules use `hit@100` at the last prompt token, exclude the last layer and require at least 10 development records per category. *Rule v1* (Qwen3.5-4B, Gemma-4-E4B-it). Band: layers where correctly answered round-0 records read the original bridge at ≥ 0.5 above control under the J-lens. Categories: band mean ≥ 0.5 . Validation (F2, “adoption”): the peer bridge under the hypothetical condition exceeds the peer bridge under mention-bridge by ≥ 0.3 , with ≥ 30 records, and the original bridge in agree dialogues is at least half its round-0 value. *Rule v2* (Qwen3.6-27B, Llama-3.1-8B-Instruct). Categories: over the round-0 window, the span from the first to the last layer where the J-lens reads the original bridge at least 0.25 above the logit lens, categories whose window mean of that difference is ≥ 0.25 . Band: layers where agree records read the original bridge at ≥ 0.3 under the J-lens and at least 0.25 above the logit lens. Validation: hypothetical $\geq 3\times$ mention-bridge, with at least 30 records and the same persistence check.

Development-stage adjustments, as recorded in the files. *Qwen3.5-4B.* The declared threshold `hit@10` gave an empty band in both formats (the bridge is rarely in the lens top 10 at the answer position; J-lens peak 0.48), so `hit@100` was adopted (the control’s `hit@100` in the band is 0.02–0.05). The last layer was excluded, since it is the model output and identical for every lens. The declared retention check (F2) could not be evaluated for lack of sibling facts in readable categories and was replaced by the adoption check. Items were restricted to categories whose bridge is readable in round 0. *Qwen3.6-27B.* The copied rule v1 selected only layer 62 of 64, where both lenses read the bridge (development round 0: J-lens 0.722, logit lens 0.717), so H4 could not be tested; the file was kept and its test split never run. Rule v2 was written on this model’s development split before any of its test data. The adoption check was changed from a gap to a ratio because readout levels differ between models; under the gap, this model’s development data would have failed ($0.331 - 0.053 = 0.278 < 0.3$; the ratio is 6.2). *Llama-3.1-8B-Instruct.* Rule v1 failed in both formats (J-lens peak

*Both authors contributed equally to this research.

Table 1: Registration files, with the time each was first committed (UTC) and the first 16 hex digits of its SHA-256. Each file is included verbatim in the repository. The test queues refuse to start unless the model’s registration is committed unchanged.

File	Registers	First committed (UTC)	SHA-256 (prefix)
qwen35_4b.json	letter study: readout layer 23 of 32	2026-09-26 00:42	08915272776af11d
qwen35_4b_addendum1.json	letter study: positive control (instructed)	2026-09-26 17:08	d8ab80b983b06d96
qwen35_4b_entity.json	entity study, Qwen3.5-4B: rule v1, layers 23–30, H1–H4	2026-09-27 02:05	2a85d61dce175353
qwen35_4b_entity_rounds.json	later rounds (summary protocol): R1, R2	2026-09-28 11:14	dd128b94eadeb1b5
qwen36_27b_entity.json	Qwen3.6-27B, rule v1 (layer 62); superseded, test never run	2026-09-29 14:39	f46838032349df8a
qwen36_27b_entity_v2.json	Qwen3.6-27B, rule v2, layers 52–55	2026-09-29 15:24	f9b61ee2afdc34ec
llama31_8b_it_entity.json	Llama-3.1-8B-Instruct, rule v2, chat, layers 17–23	2026-09-29 15:47	e7456aba8bae0cf3
gemma4_e4b_it_entity.json	Gemma-4-E4B-it, rule v1, layers 32, 34–40	2026-09-29 21:02	13b1d20a2cb5c70c
qwen35_4b_entity_own.json	own-answer addendum (hidden, absent): O1–O4	2026-09-30 03:34	26378d0b12c05169
qwen36_27b_entity_v2_own.json	own-answer addendum, Qwen3.6-27B	2026-09-30 03:34	c8b0b2757f87d997
llama31_8b_it_entity_own.json	own-answer addendum, Llama-3.1-8B-Instruct	2026-09-30 03:34	e04e7a03e96efaa1
gemma4_e4b_it_entity_own.json	own-answer addendum, Gemma-4-E4B-it	2026-09-30 03:34	bb69b17c0392f01e

0.30 at layer 24 in the plain format, where the logit lens reads as much), so no v1 registration was written. Rule v2, taken unchanged from Qwen3.6-27B, selected nothing in the plain format and passed in the chat format. *Gemma-4-E4B-it*. Rule v2 selected nothing in either format: the bridge becomes readable at layers 30–31 under every lens at once. Rule v1 was applied unchanged, and the registration states that H4 is expected to be null on this model. The primary lens is a refit on the instruction-tuned model (Section 5). *All replications* run round 1 only; later rounds were registered only for Qwen3.5-4B.

Letter-study addendum. The positive-control addendum records that it was written after the main and control pressure runs of the letter study had been analysed, including exploratory analyses at every layer; it adds one condition and fixes its analysis and decision rule in advance, without changing the readout layer or definitions.

3 Prompts

All prompts are built by `silent_dissent/entity_prompts.py`; the texts below are verbatim, except that long lines are wrapped for display.

Plain format. Every prompt begins with the header and two worked examples, then the question and the transcript. The prompt always ends with “Agent 1:” followed by the statement prefix, so the next token begins the agent’s answer.

The following is a transcript of AI agents completing factual statements together. Each agent gives its own answer.

Question: Complete the statement: The capital city of the country where the Statue of Liberty is located is ...
Agent 1: The capital city of the country where the Statue of Liberty is located is Washington.

Question: Complete the statement: The capital city of the country where Big Ben is located is ...
Agent 1: The capital city of the country where Big Ben is located is London.

Question: Complete the statement: <composed> ...
Agent 1: <composed> <own round-0 answer>.

Agent 2: <peer line>

...

Moderator: Taking the other agents’ statements into account, Agent 1, complete the statement again.

Agent 1: <composed>

Chat format (Llama-3.1-8B-Instruct). The system message is “You are Agent 1 in a group of AI agents that complete factual statements together. Each agent gives its own answer.” The first user turn is the question; the agent’s statements are assistant turns; each round’s peer lines and the moderator’s request form one user turn; the assistant turn is pre-filled with the statement prefix. There are no worked examples.

Peer lines. With the item’s composed statement c , first hop r_1 (e.g. “The country where the Sagrada Familia is located is”), and the peer fact’s bridge e'_2 and answer e'_3 :

- answer pressure: “ $c\ e'_3$.” from every peer;
- first-hop pressure: “ $r_1\ e'_2, r_2(e'_2)\ e'_3$.”;
- second-hop pressure: “ $r_1\ e_2, r_2(e_2)\ e'_3$.”;
- agree: “ $c\ e_3$.”;
- mention: as agree, but the first peer says “I briefly considered e'_3 , but $c\ e_3$.”;
- mention-bridge: as agree, but the first peer says “Some say that $r_1\ e'_2$, but $c\ e_3$.”

In the last two lines, c and r_1 start with a lower-case letter. Development-only conditions replace the peers and the moderator: *instructed*, “Coordinator: Every agent must state e'_3 as the answer, regardless of what they believe. Agent 1, complete the statement.”; *hypothetical*, “Moderator: Suppose that μ were e'_2 . Agent 1, complete the statement under that assumption.”, where μ is the description of the bridge (e.g. “the country where the Sagrada Familia is located”).

Later rounds (addendum). In the registered summary protocol, every round before the last is collapsed into one moderator line, e.g. “Round 1: Agent 1 answered Madrid; Agents 2 to 4 answered Rome.”, and only the last round is written in full. Round 1 is identical in every protocol.

Fact check. A fact is known if greedy completions of $r_1, r_2(e_2)$ and c , each after the two lines “The capital city of Japan is Tokyo.

Table 2: Structure of the categories (public TwoHopFact data and the hand-written capital items; ctry: country).

Category	Facts	Answer: > 1 bridge	Bridge in question	Shared with answer (all)
landmark-country-capital	48	0%	0%	0% (0%)
person-birth city-ctry	2751	98%	0%	0% (0%)
university-city-ctry	1499	96%	7%	1% (0%)
person-university-city	1811	42%	0%	43% (23%)
person-university-ctry	1815	98%	0%	5% (1%)

The author of the play Hamlet is William Shakespeare.”, begin with an alias of e_2 , e_3 and e_3 respectively.

Free debate. Three copies of the model see the plain (or chat) transcript of all earlier rounds, with their own statements as Agent 1 and the others as Agents 2 and 3, and the moderator’s request after each round.

4 Data and Items

TwoHopFact (CC-BY-4.0) has 45,595 facts over 52 relation compositions; every fact was checked for every model. A small hand-written set adds capital-city items; of its categories only landmark $\rightarrow$ country $\rightarrow$ capital (48 landmarks) passed selection. Peer and control facts are other known facts of the same category and split, drawn at random per item, with bridges and answers distinct from the item’s and from each other; items without two such partners are dropped.

Table 2 describes the structure of the categories the study uses, computed on the public data by `scripts/twohopfact_audit.py` (all facts of each category, not only those a model knows). Only in the capital items does the peers’ answer determine their bridge. In the university-city category, some university names contain their city, so for those items the bridge is partly stated in the question. In the person-university-city category the bridge often shares words with the answer (e.g. a university named after its city); shared tokens are removed from the probes, and items left with no bridge probe drop out of the statistics.

5 Readout and Lens Details

Probes. An entity’s probe tokens are the first tokens of “w” for each of its words w with at least three letters that is not a stop word (the list, in `entity_facts.py`, includes articles and particles, *city*, *university*, *college*, company suffixes, *republic*, *kingdom*, *new* and the compass directions). Tokens that occur in two or more of an item’s six roles (original, peer and control bridge and answer) are removed from all of them; an entity without a content word falls back to the first token of its name. The rank of a role is the minimum full-vocabulary rank of its remaining tokens; ranks are capped, and a role with no token is missing.

Positions. Ranks are stored for every layer at two positions, the last token of the subject e_1 and the last prompt token; the paper uses the last prompt token. The development runs of the person-bridge follow-up also stored the mention of the bridge’s description and the minimum over the span from the description to the last token.

Lenses. The logit lens applies the model’s final norm and unembedding. The J-lens first multiplies by J_ℓ , the Jacobian of the final residual with respect to the layer- ℓ residual, summed over target positions at or after each source position, then averaged over source positions and sequences, with the first 16 positions (attention sinks) excluded, on 128-token WikiText sequences; our estimator is a port of the reference implementation, is tested against a brute-force Jacobian and matches the reference code to 10^{-8} on a shared model. The Qwen and Llama lenses are the pre-fitted ones hosted by Neuronpedia in the Hugging Face repository `neuronpedia/jacobian-lens`, branch `qwen-n1000` (commit `16a01f3`, unchanged since June 2026): `Qwen3.5-4B_jacobian_lens_n1000.pt`, `Qwen3.6-27B_jacobian_lens_n1000.pt` ($d_{\text{model}} = 5120$, one lens per layer below the last of 64) and `Llama-3.1-8B-Instruct_jacobian_lens.pt`; the paths are in the configuration files. For Gemma-4-E4B-it the hosted lens (`gemma-4-E4B_jacobian_lens.pt`, branch `main`, commit `b25d72a`) was fitted on the base model; the primary lens is a refit on the instruction-tuned model with 100 sequences, and the base-model lens is a secondary readout. Gemma’s output logits are soft-capped ($30 \tanh(x/30)$); the lens applies the same cap, which does not change ranks. On Qwen3.5-4B and Llama-3.1-8B-Instruct, 100-sequence refits read like the hosted lenses (per-layer differences of about 0.01–0.05 on development data).

Injection directions. The direction of token t at layer ℓ is $J_\ell^\top (g \odot W_U[t])$, where g is the elementwise gain of the final norm, measured as $\text{norm}(1)$ so that both w and $(1+w)$ parameterizations are covered; an entity’s direction is the normalized mean over its probe tokens. The random direction is a Gaussian vector, fixed per item and layer, normalized to unit length. The injection adds $\alpha \|h\| d$ to the last position of one layer.

Statistics. $\text{Hit}@k$ is 1 if a role’s rank is below k . A record’s readout is the band mean of $\text{hit}@k(\text{role}) - \text{hit}@k(\text{control bridge})$. Confidence intervals are percentile bootstrap intervals over records with 1,000 resamples (one-sample and paired statistics) or 2,000 (two-sample differences), seed 0.

6 Additional Results

Conformity. Table 3 gives the share of agents stating the peers’ answer for every kind of pressure and number of peers, and the share of other answers. The ordering of pressure kinds reported in the paper holds for three peers. With five peers, first-hop pressure persuades less than answer pressure for Qwen3.5-4B and Qwen3.6-27B, and second-hop pressure persuades more than answer pressure for Gemma; with one peer, where all rates under answer and second-hop pressure are tiny, second-hop pressure persuades slightly more than answer pressure for Qwen3.6-27B and Gemma.

Conformity and retention. Figure 1 plots the share of agents that gave in against the ratio of the J-lens readout of agents that gave in to that of agents whose peers agreed (paper §4.7), at the registered $k = 100$ and at $k = 500$. Only the two ends of the range line up with conformity, the ratio compares different item sets and it has no interval; the paper reports it as an observation.

Readouts per layer. The per-layer readouts are Figure 3 of the paper (§4.6).

Table 3: Share (%) of agents giving the peers’ answer in round 1, by kind of pressure and number of peers (top), and share (%) of other answers (bottom). n : items entering round 1.

Model	n	answer			hop 1			hop 2		
		1	3	5	1	3	5	1	3	5
Qwen3.5-4B	668	2	8	8	10	8	5	0	0	0
Qwen3.6-27B	476	1	39	49	26	49	43	1	6	6
Gemma-4-E4B-it	476	0	25	8	37	52	29	1	24	9
Llama-3.1-8B-Instruct	542	10	90	89	80	97	95	3	8	6
Qwen3.5-4B (other)		0.0	0.0	0.4	0.0	0.3	0.1	0.0	0.0	0.0
Qwen3.6-27B (other)		0.0	0.2	0.2	0.0	0.0	0.0	0.0	0.0	0.0
Gemma-4-E4B-it (other)		0.2	0.4	1.7	1.3	3.1	3.4	0.0	0.8	0.8
Llama-3.1-8B-Instruct (other)		1.3	3.1	2.8	0.5	0.4	0.7	2.0	1.1	1.7




Figure 1: Conformity against retention (descriptive): share of agents that gave in (answer pressure, three peers) and the J-lens readout of agents that gave in divided by that of agents whose peers agreed, at $k = 100$ and $k = 500$. Numerator and denominator are different item sets and the ratio has no uncertainty estimate.

Readouts per condition. Tables 4–7 give the band-mean readouts of the original and the peer bridge for every condition and lens, with 95% intervals. “Held” and “gave in” are the agents under answer pressure with three peers that kept their answer or stated the peers’ answer; the mention conditions include only agents that kept their answer.

Robustness. Table 8 lists H1–H4 under the other band rule (where it selects layers), other rank thresholds, one and five peers, and the secondary lenses. All rows are exploratory. Figure 2 shows the readouts of agents that gave in and of agents whose peers agreed as a function of k .

Development-split validation. Table 9 gives the validation checks recorded in each registration (development split, the registered format). The instructed condition, which the paper discusses (paper §6), reads the original bridge near zero in three models; the hypothetical condition, in which the agent adopts the peer bridge on request, reads it well above the mention-bridge condition in every model.

Injection (E4). Table 10 gives the share of agents whose most likely next token is the first token of the original answer (agents that gave in) or of the peers’ answer (agents that held), for each direction and strength, pooled over the three layers. In agents that held, the peers’ answer direction moves 87–89% of them to the peers’ answer in every model (maximum over α); the peer bridge direction exceeds the control bridge in Qwen3.5-4B (0.12 vs. 0.02), Qwen3.6-27B (0.13 vs. 0.03) and Llama (0.20 vs. 0.10), but not in Gemma (0.01), while random directions reach 0.03–0.10.

Later rounds (Qwen3.5-4B addendum). On 200 development items, the registered repeat protocol produced “other” answers (e.g. “correct”, “identical”) in up to 58% of a cell’s agents by round 3, a rephrased protocol in up to 78%, and the summary protocol in at most 1.1%, with round-1 answers identical to the reference in all 1,050 agents; the summary protocol was registered. On the test split, $R1 = 0.541$ [0.469, 0.623] ($n = 49$) and $R2 = +0.409$ [0.291, 0.540]; the change of the R1 statistic from round 1 to round 3 in the same agents is -0.324 [−0.383, −0.258]. Exploratory: in the later-round run, the share of agents giving in under first-hop pressure rose from 8.5% in round 1 to 29% in round 2.

Timeline (E3, Qwen3.5-4B). Reading every token of the final round shows the peers’ bridge building up as each successive peer speaks. The timeline was also queued for the replications, but only the Qwen3.5-4B run was analysed.

Free debate (E5). Table 11 gives the group accuracy per round. Round-0 answers are sampled at temperature 1 and all agents share the round-0 prompt, so round-0 latent votes are one readout counted three times and are not comparable with stated votes. Qwen3.6-27B’s stated accuracy drops in rounds 2 and 3 while its latent votes do not.

Answers in free debate (exploratory; rule recorded before the answers were read). Every answer was classed as correct, another round-0 candidate, degenerate (empty; first word a verdict such as *correct*, *incorrect*, *identical*, *same*, *agree(d)*, *right*, *true*, *yes*, *no* or *unchanged*; starting with “the same”; or containing the word “agent”) or other (Table 12). The drop was to be attributed to the protocol if degenerate answers made up at least 10% of round 2 or 3 and if, with them treated as abstentions, the plurality accuracy (b) at both rounds was within 0.05 of round 1 or its interval overlapped




Figure 2: Original-bridge readout (band mean, $\text{hit}@k$ minus control) as a function of the rank threshold k from 1 to 10,000, for agents that gave in (solid) and agents whose peers agreed (dashed), J-lens and logit lens (exploratory; Section 8). The registered threshold is $k = 100$. The J-lens readout falls at large k because the control entity then also reaches the top k ; most logit-lens curves still rise at $k = 10,000$.

round 1’s (2,000 bootstrap resamples over facts). Neither held for Qwen3.6-27B (degenerate answers at most 8%; (b) 0.31 and 0.46 against 0.78), and no model met the rule, so the drop is not attributed. At round 2, 51% of Qwen3.6-27B’s answers were “other”: strings no agent had given in round 0, the most frequent being “an island” (45), “an island nation in East Asia” (27), “agents” (25), “the agents” (24), “loop” (11) and “unconfirmed” (9). Of the 648 agents that were right at round 1, 414 switched at round 2 to an answer no peer had given at round 1 and 2 adopted a peer’s wrong answer (labels as defined below): the drop is not conformity. Accuracy (a) reproduces Table 11 exactly (the script checks it). The rule says “the word agent”, and the implementation matches that word in the singular only; its description (“a mention of an agent”) reads more broadly.

Post hoc variants (not fixed in advance). After reading this result, we recomputed the rule with broader abstentions (Table 13). Adding answers with a word beginning with “agent” leaves the verdict unchanged. Treating every answer that is neither correct nor a round-0 candidate as an abstention, an upper bound that also drops genuine wrong entities, attributes the drop, but leaves 139 of 279 facts at round 2 without an answer. No other model is attributed under any variant. We therefore report the drop as unexplained.

Original bridges in free debate (exploratory; rules recorded before the run). Each agent’s prompt in every round was rebuilt from the stored transcripts exactly as the debate had built it; nothing was generated. Before any readout, the rebuilt round-1 prompts had to reproduce the stored J-lens ranks of the candidate answers (at least 95% within 0.1 in $\log_{10}(\text{rank} + 1)$), and they reproduced all of them exactly in every model. At the last token we read the original bridge and a control bridge (a fact of the same category with a different bridge and answer, drawn with the configuration’s seed) with each model’s registered lens, band and k , and with the logit lens. Agents right at round $r - 1$ were labelled *gave in* (wrong at r , stating a wrong answer another agent gave at $r - 1$), *switched* (wrong, not a peer’s answer), *held* (right, although a peer gave a

wrong, non-degenerate answer at $r - 1$), *held degenerate* (right, the only wrong peers having given degenerate answers) or *agreed* (all peers right). Intervals resample facts (10,000 resamples, seed 0) and are given only for cells with at least 20 readable records. Known facts at round 1 are primary, and silent dissent in free debate was to be claimed only if the J-lens readout of agents that gave in had an interval above zero with at least 20 records. Only Llama had enough such agents (Table 14): 0.20 [0.11, 0.32] ($n = 29$, 22 facts), J-lens minus logit lens 0.18 [0.08, 0.29], and $-0.03 [-0.19, 0.13]$ relative to agents that held. On all facts, a secondary analysis, Qwen3.5-4B’s 20 agents that gave in read 0.64 [0.41, 0.85]. The comparison with the main study is descriptive: in free debate the two peers are copies of the same model that may disagree with each other, and which agents give in is not controlled.

Lens fitting corpus (E6, Qwen3.5-4B). Lenses refitted on dialogue text, on a mixed corpus and on 100 WikiText sequences change token-level fidelity, but change the entity readouts of the development conditions by about 0.02 or less; 100 sequences suffice.

7 Own-Answer Addendum

Design. The addendum asks whether agents that gave in read their original bridge because their own round-0 answer is in context. It repeats round 1 (three peers; answer and first-hop pressure; agree and mention baselines) with the agent’s round-0 statement replaced by “(answer submitted)” (*hidden*) or left out (*absent*). Under absent the agent answers for the first time after the peers, and the moderator asks it to “complete the statement” instead of “...again”; in the chat format the question, the peers and the moderator then form one user turn. Round 0 is still run alone to select the items answered correctly. The variants, cells, pilot size and checks were declared in each model’s configuration before any run.

Development pilots and registration. On up to 300 development items per variant, a variant passed if every cell had at most 10% “other” answers and agents whose peers agreed read the original bridge with $n \geq 30$ and a lower interval bound above zero (Table 15).

Both variants passed for every model and were registered; readout, format and categories are those of each model’s registration, nothing was selected, and a hypothesis whose test cell has fewer than 10 records gets no verdict. The test runs started after the addenda were committed.

Hypotheses and results. For agents that gave in under answer pressure: O1, the original bridge reads above control under hidden; O3, under absent; O2, higher under hidden than under absent; O4, higher under the J-lens than under the logit lens under hidden (paired). The verdicts are in the paper (Table 3): O1, O3 and O4 are supported for every model, O2 only for Gemma. Hidden and absent differ in turn structure and wording as well as in whether the agent answered before, and O2 compares different agents, so O2 is not a clean test of commitment; we recorded this caveat before the test runs. Table 16 gives the give-in rates per variant (exploratory), the per-item comparisons of agents that gave in under two variants, and the peer-bridge comparison with the mention baseline; Table 17 the readouts per condition. For Llama, the readout of agents whose peers agreed also drops when the earlier answer is hidden (0.24) or absent (0.16), against 0.52 with it visible; under absent its agents that gave in read more than agents whose peers agreed (0.39 vs. 0.16), under hidden about as much (0.25 vs. 0.24).

8 Exploratory Robustness Analyses

These analyses answer questions raised after the test runs. Their choices were recorded before they were run: the registered readout of each model; bridge-clustered intervals with $B = 10,000$ resamples (seed 0), clustering on the original bridge for H1, H3 and H4 and on the peers’ bridge for H2, one draw shared by both samples of a difference; two-sided bootstrap p -values; a Holm correction over every registered test with an estimate (15), family-wise $\alpha = 0.05$; intervals per category only where each sample has at least 10 readable records. None of them changes a registered verdict.

Clustered intervals and multiple tests. Table 18. Fourteen of the 15 tests remain supported after clustering and correction; Llama’s H2 does not. With records resampled independently, all 15 remain supported after the correction.

Without a threshold. For each record and band layer, logrank is $\log_{10}(r_c + 1) - \log_{10}(r_b + 1)$ for the ranks r_b of the bridge and r_c of the control (clipped at 30,000), and win is 1 if the bridge outranks the control, 1/2 on a tie and 0 otherwise; both are averaged over the band (Table 19). Under the J-lens, agents that gave in rank their original bridge above the control in every model. The J-lens-logit difference persists for Qwen3.5-4B and Llama, vanishes for Qwen3.6-27B and reverses for Gemma, where the logit lens ranks the bridge above the control about as much or slightly more, mostly outside the top 100 (Figure 2); Llama’s logit lens ranks it above the control only marginally.

Same items. Table 20 compares each item with itself: the readout of an agent that gave in against the same item’s agree, round-0 or mention readout, with the ratio of the two means for the registered readout. Giving in lowers the original-bridge readout in every model. The selection rows compare items whose agent gave in with items whose agent held, before any pressure: in Qwen3.5-4B

and Gemma the items that gave in read higher in round 0 and when peers agreed, which is consistent with Qwen3.5-4B’s agents that gave in reading more than agents whose peers agreed across item sets.

Categories. Table 21. No agent of three models gave in on capital items, the only category in which the peers’ answer determines their bridge; the 14 agents of Qwen3.6-27B that did (of 26) read the peers’ bridge only 0.02 $[-0.10, 0.15]$ above the mention baseline. H2 therefore rests on categories in which the peers’ bridge is associated with their answer. H1 is positive in every category with at least 10 readable records.

9 Negative Results

Letter-answer study (Qwen3.5-4B). Four-option multiple-choice items (MMLU and ARC) were debated with scripted peers, and the letter logits were read at the last prompt token after the forced prefix “The answer is”. The registered readout layer (23 of 32) was the earliest at which the J-lens top letter matched the model’s output on at least 90% of round-0 items. Strict silent dissent (the agent changed its letter while its original letter was still top among the letters at layer 23) occurred in about 6% of round-1 flips and about 1.6% by round 3; injecting the original letter’s direction restored the original answer in 10–14% of cases. The registered positive control (instructed compliance, addendum) asked whether the readout detects a known split: a coordinator required every agent to submit a given wrong letter regardless of its belief. Compliance was 100%, and strict silent dissent was 0 of 1,000, below the registered “blind” threshold (upper confidence bound < 0.2). The letter readout therefore cannot distinguish internalized flips from hidden dissent, and the low rate under pressure is not evidence of sincere persuasion. Two exploratory checks qualify this verdict: over the last 16 prompt positions and all 32 layers, no position was both reliable and sensitive; and when asked to report the letter that follows the one it believes correct (e.g. $D \rightarrow A$), with its own round-0 answer visible, the logit lens at position -2 , layer 20, read the model’s own answer in 74% of records (here the logit lens did better than the J-lens), so a lens can see an answer that is computed and then transformed. Outside the letter task, the J-lens was better than the logit lens on generic text (mean KL divergence to the output over layers 8–24: 6.5 vs. 9.5 on WikiText) and on unstated two-hop bridges (peak readout 0.79 vs. 0.30, readable from about layer 7), which motivated the entity study.

Person bridges (Qwen3.5-4B). Bridges that are people (e.g. the author of a novel) were not readable at the answer position. A planned addendum declared its rules in the configuration before its development runs: readouts at several positions and probe modes (all content words or the surname), selected by the main registration’s thresholds. Person names were readable when the name was the next word (positive control: readable from about layer 23, up to 0.91), but in two-hop statements only the minimum over the span from the description to the answer position reached 0.58 (layer 30), and the adoption check had 20 records against the declared minimum of 30 (adoption 1.0 vs. 0.16). The addendum failed its own rules and was neither registered nor run on the test split; we did not relax the rules.

Own-answer addendum and exploratory robustness. O2 (a higher readout under hidden than under absent) failed for three of the four models, with the difference below zero for Qwen3.5-4B and Llama (Section 7). Among the exploratory analyses, Llama’s H2 is no longer supported with bridge-clustered intervals and a Holm correction, the J-lens–logit difference does not hold without a threshold for Qwen3.6-27B and Gemma, and the peers’ bridge is not read in the one category where their answer implies it (Section 8).

10 Compute

All runs used one NVIDIA A100 80GB. The test-split queue of the Qwen3.5-4B entity study together with the lens-corpus experiments took about 4 hours, its follow-ups about 1 hour, and the three replications about 11.2 hours of wall time, including model downloads, fact checks over all 45,595 facts, lens refits, development runs, test runs and analyses. The own-answer addendum (development pilots and test runs for all four models) took about 1.2 hours and the exploratory robustness analyses a few minutes on the same machine. The letter study and the Qwen3.5-4B development runs are not

included in these figures. Figures and tables are produced from the summary file on a CPU in seconds.

11 Use of AI Tools

The authors worked with Claude Code (Anthropic; model Claude Opus 5.5) in interactive sessions throughout the project. It wrote and ran most of the experiment and analysis code under the authors’ direction (including the GPU job queues), proposed and discussed design choices that the authors then decided on, among them the entity-answer paradigm after the letter study’s positive control came out blind, the layer-selection rules v1 and v2, the control conditions and the later-round protocols, drafted documentation, and edited the text of the paper. There is no single prompt: the work proceeded as a sequence of conversational sessions, whose logs the authors keep and will share on request. Every pre-registration was committed before the corresponding test data existed, as recorded in Table 1, and every number in the paper was checked against the summary data, the registration files and the experiment records.

Table 4: Qwen3.5-4B: band-mean readouts (hit@100 minus control), 95% CI.

Condition	original bridge		peer bridge	
	J-lens	logit	J-lens	logit
round 0 alone ($n=663$)	0.69 [0.66, 0.72]	0.26 [0.24, 0.29]	-0.01 [-0.02, 0.01]	0.00 [0.00, 0.01]
peers agree ($n=663$)	0.63 [0.60, 0.66]	0.19 [0.18, 0.21]	0.00 [-0.01, 0.01]	0.00 [0.00, 0.00]
mention ($n=660$)	0.62 [0.58, 0.65]	0.17 [0.15, 0.18]	0.03 [0.01, 0.04]	0.00 [0.00, 0.00]
mention-bridge ($n=661$)	0.62 [0.59, 0.65]	0.15 [0.13, 0.17]	0.18 [0.16, 0.19]	0.03 [0.02, 0.03]
held (answer) ($n=609$)	0.54 [0.50, 0.57]	0.09 [0.08, 0.10]	0.16 [0.14, 0.19]	0.00 [0.00, 0.00]
gave in (answer) ($n=54$)	0.85 [0.78, 0.92]	0.06 [0.04, 0.09]	0.40 [0.29, 0.51]	0.01 [0.00, 0.03]

Table 5: Qwen3.6-27B: band-mean readouts (hit@100 minus control), 95% CI.

Condition	original bridge		peer bridge	
	J-lens	logit	J-lens	logit
round 0 alone ($n=449$)	0.63 [0.59, 0.68]	0.07 [0.06, 0.09]	0.00 [-0.01, 0.01]	0.00 [0.00, 0.01]
peers agree ($n=448$)	0.50 [0.45, 0.54]	0.07 [0.06, 0.09]	-0.01 [-0.02, 0.00]	0.00 [0.00, 0.01]
mention ($n=445$)	0.43 [0.39, 0.48]	0.02 [0.01, 0.02]	0.03 [0.01, 0.06]	0.00 [0.00, 0.01]
mention-bridge ($n=385$)	0.19 [0.16, 0.23]	0.02 [0.01, 0.03]	0.06 [0.04, 0.08]	0.01 [0.00, 0.01]
held (answer) ($n=280$)	0.40 [0.35, 0.45]	0.01 [0.00, 0.02]	0.03 [0.01, 0.04]	0.00 [0.00, 0.01]
gave in (answer) ($n=168$)	0.22 [0.17, 0.28]	0.00 [0.00, 0.00]	0.16 [0.12, 0.20]	0.00 [0.00, 0.00]

Table 6: Gemma-4-E4B-it: band-mean readouts (hit@100 minus control), 95% CI. J-lens (base): the lens hosted by Neuronpedia, fitted on the base model.

Condition	original bridge			peer bridge		
	J-lens	logit	J-lens (base)	J-lens	logit	J-lens (base)
round 0 alone ($n=471$)	0.82 [0.79, 0.85]	0.78 [0.75, 0.81]	0.80 [0.77, 0.83]	0.00 [-0.01, 0.01]	0.00 [-0.01, 0.01]	0.00 [-0.01, 0.01]
peers agree ($n=471$)	0.67 [0.63, 0.70]	0.57 [0.53, 0.60]	0.55 [0.51, 0.58]	0.00 [-0.01, 0.01]	0.00 [-0.01, 0.01]	0.00 [-0.01, 0.01]
mention ($n=471$)	0.65 [0.62, 0.69]	0.53 [0.50, 0.57]	0.52 [0.49, 0.56]	0.00 [-0.01, 0.01]	0.00 [-0.01, 0.00]	0.00 [-0.01, 0.00]
mention-bridge ($n=470$)	0.77 [0.74, 0.80]	0.70 [0.66, 0.73]	0.72 [0.69, 0.75]	0.01 [-0.01, 0.02]	0.00 [-0.01, 0.01]	0.00 [-0.01, 0.01]
held (answer) ($n=349$)	0.61 [0.57, 0.66]	0.54 [0.50, 0.58]	0.55 [0.51, 0.59]	0.02 [0.00, 0.03]	0.00 [0.00, 0.01]	0.00 [-0.01, 0.01]
gave in (answer) ($n=120$)	0.24 [0.17, 0.30]	0.05 [0.03, 0.07]	0.11 [0.06, 0.15]	0.21 [0.17, 0.26]	0.06 [0.03, 0.08]	0.02 [0.01, 0.04]

Table 7: Llama-3.1-8B-Instruct: band-mean readouts (hit@100 minus control), 95% CI. J-lens (refit): our 100-sequence refit.

Condition	original bridge			peer bridge		
	J-lens	logit	J-lens (refit)	J-lens	logit	J-lens (refit)
round 0 alone ($n=538$)	0.48 [0.45, 0.52]	0.12 [0.10, 0.14]	0.47 [0.43, 0.50]	0.00 [0.00, 0.01]	0.00 [0.00, 0.00]	0.00 [-0.01, 0.01]
peers agree ($n=535$)	0.52 [0.48, 0.55]	0.07 [0.06, 0.09]	0.47 [0.44, 0.51]	0.00 [-0.01, 0.01]	0.00 [0.00, 0.00]	0.00 [-0.01, 0.01]
mention ($n=538$)	0.47 [0.43, 0.50]	0.09 [0.07, 0.11]	0.43 [0.40, 0.47]	0.01 [0.00, 0.01]	0.00 [0.00, 0.00]	0.00 [0.00, 0.01]
mention-bridge ($n=537$)	0.20 [0.17, 0.23]	0.04 [0.03, 0.05]	0.17 [0.14, 0.20]	0.07 [0.05, 0.09]	0.01 [0.01, 0.02]	0.06 [0.04, 0.08]
held (answer) ($n=36$)	0.29 [0.17, 0.41]	0.04 [0.01, 0.08]	0.24 [0.14, 0.36]	0.01 [-0.04, 0.05]	0.00 [0.00, 0.00]	0.00 [-0.04, 0.03]
gave in (answer) ($n=485$)	0.03 [0.02, 0.05]	0.00 [0.00, 0.01]	0.02 [0.01, 0.04]	0.03 [0.02, 0.04]	0.02 [0.01, 0.03]	0.02 [0.01, 0.03]

Table 8: Robustness on the test split (exploratory): estimate and 95% bootstrap interval. n : records for H1.

Model	Variant	H1	H2	H3	H4
Qwen3.5-4B	band of rule v2 ($n=54$)	0.818 [0.735, 0.895]	0.362 [0.251, 0.483]	–	0.812 [0.728, 0.886]
Qwen3.5-4B	$k = 10$ ($n=54$)	0.123 [0.081, 0.167]	0.043 [0.009, 0.082]	–	0.120 [0.079, 0.167]
Qwen3.5-4B	$k = 50$ ($n=54$)	0.643 [0.551, 0.732]	0.226 [0.141, 0.320]	–	0.627 [0.539, 0.715]
Qwen3.5-4B	$k = 500$ ($n=54$)	0.701 [0.595, 0.806]	0.373 [0.212, 0.528]	–	0.546 [0.447, 0.653]
Qwen3.5-4B	1 peers ($n=11$)	0.784 [0.636, 0.932]	0.266 [0.079, 0.473]	0.260 [0.184, 0.340]	0.682 [0.545, 0.841]
Qwen3.5-4B	5 peers ($n=53$)	0.816 [0.734, 0.894]	0.392 [0.280, 0.507]	0.379 [0.270, 0.488]	0.755 [0.670, 0.837]
Qwen3.6-27B	band of rule v1 ($n=168$)	0.726 [0.655, 0.792]	0.335 [0.250, 0.419]	0.194 [0.008, 0.319]	0.059 [-0.012, 0.131]
Qwen3.6-27B	$k = 10$ ($n=168$)	0.002 [0.000, 0.004]	-0.001 [-0.002, 0.000]	0.060 [-0.004, 0.146]	0.002 [0.000, 0.004]
Qwen3.6-27B	$k = 50$ ($n=168$)	0.048 [0.028, 0.070]	0.023 [0.005, 0.041]	0.197 [0.058, 0.359]	0.048 [0.028, 0.070]
Qwen3.6-27B	$k = 500$ ($n=168$)	0.414 [0.319, 0.513]	0.214 [0.111, 0.316]	-0.099 [-0.485, 0.267]	0.384 [0.290, 0.485]
Qwen3.6-27B	1 peers ($n=1$)	-0.250 [-0.250, -0.250]	–	0.225 [-0.084, 0.549]	-0.250 [-0.250, -0.250]
Qwen3.6-27B	5 peers ($n=211$)	0.235 [0.187, 0.283]	0.113 [0.080, 0.151]	0.289 [0.086, 0.513]	0.228 [0.182, 0.274]
Gemma-4-E4B-it	$k = 10$ ($n=120$)	0.000 [0.000, 0.000]	0.000 [-0.001, 0.001]	0.021 [0.006, 0.041]	0.000 [0.000, 0.000]
Gemma-4-E4B-it	$k = 50$ ($n=120$)	0.054 [0.032, 0.078]	0.022 [0.005, 0.042]	0.289 [0.228, 0.354]	0.051 [0.028, 0.075]
Gemma-4-E4B-it	$k = 500$ ($n=120$)	0.556 [0.484, 0.626]	0.423 [0.326, 0.512]	0.414 [0.340, 0.488]	0.071 [0.001, 0.147]
Gemma-4-E4B-it	1 peers ($n=0$)	–	–	0.185 [-0.185, 0.802]	–
Gemma-4-E4B-it	5 peers ($n=38$)	0.306 [0.201, 0.428]	0.269 [0.165, 0.379]	0.585 [0.465, 0.701]	0.247 [0.148, 0.352]
Gemma-4-E4B-it	lens jlens_base ($n=120$)	0.105 [0.065, 0.150]	0.023 [0.006, 0.046]	0.329 [0.259, 0.401]	0.059 [0.025, 0.099]
Llama-3.1-8B-Instruct	$k = 10$ ($n=485$)	0.002 [0.000, 0.004]	-0.000 [-0.002, 0.001]	0.029 [0.006, 0.062]	0.001 [0.000, 0.002]
Llama-3.1-8B-Instruct	$k = 50$ ($n=485$)	0.014 [0.006, 0.023]	0.005 [-0.002, 0.014]	0.292 [0.176, 0.409]	0.011 [0.005, 0.019]
Llama-3.1-8B-Instruct	$k = 500$ ($n=485$)	0.214 [0.180, 0.249]	0.316 [0.269, 0.367]	0.416 [0.231, 0.572]	0.193 [0.162, 0.226]
Llama-3.1-8B-Instruct	1 peers ($n=52$)	0.082 [0.000, 0.179]	0.091 [0.022, 0.166]	0.066 [-0.126, 0.275]	0.060 [-0.008, 0.143]
Llama-3.1-8B-Instruct	5 peers ($n=479$)	0.074 [0.055, 0.095]	0.068 [0.050, 0.086]	0.394 [0.235, 0.543]	0.063 [0.047, 0.081]
Llama-3.1-8B-Instruct	lens jlens_wiki100 ($n=485$)	0.024 [0.013, 0.037]	0.016 [0.005, 0.028]	0.358 [0.219, 0.492]	0.021 [0.010, 0.032]

Table 9: Development-split checks from the registrations: band-mean readouts (hit@ 100 minus control, 95% CI) and compliance rates. These are different items from the test split.

Check	Qwen3.5-4B	Qwen3.6-27B	Gemma-4-E4B-it	Llama-3.1-8B-Instruct
format	plain	plain	plain	chat
original bridge, round 0	0.73 [0.69, 0.77]	0.62 [0.57, 0.68]	0.79 [0.75, 0.83]	0.48 [0.44, 0.53]
original bridge, peers agree	0.66 [0.62, 0.71]	0.45 [0.40, 0.51]	0.53 [0.49, 0.58]	0.55 [0.50, 0.59]
original bridge, instructed	0.25 [0.22, 0.28]	0.04 [0.02, 0.06]	0.01 [-0.01, 0.01]	0.00 [-0.01, 0.01]
peer bridge, hypothetical	0.71 [0.66, 0.75]	0.33 [0.29, 0.37]	0.58 [0.52, 0.64]	0.43 [0.38, 0.48]
peer bridge, mention-bridge	0.12 [0.10, 0.14]	0.05 [0.03, 0.08]	0.01 [-0.01, 0.03]	0.05 [0.03, 0.07]
peer bridge, mention	0.03 [0.01, 0.04]	0.04 [0.00, 0.07]	0.00 [-0.01, 0.01]	-0.01 [-0.01, 0.00]
compliance, instructed	1.000	1.000	1.000	0.997
compliance, hypothetical	0.782	0.990	0.992	0.984

Table 10: Injection (exploratory). Gave in: share whose top next token is the original answer; held: share whose top next token is the peers’ answer. Pooled over the listed layers; items in which any two of the three answers share a first token are excluded.

Model (layers)	Agents	Direction	$\alpha=0.0$	$\alpha=0.1$	$\alpha=0.2$	$\alpha=0.4$	$\alpha=0.8$
Qwen3.5-4B (23, 26, 29)	gave in	ctrl bridge	0.04	0.07	0.08	0.06	0.00
		orig answer	0.04	0.73	0.88	0.97	0.98
		orig bridge	0.04	0.38	0.41	0.14	0.00
		random	0.04	0.04	0.06	0.15	0.23
	held	ctrl bridge	0.01	0.02	0.02	0.02	0.01
		peer answer	0.01	0.21	0.47	0.76	0.89
		peer bridge	0.01	0.06	0.11	0.12	0.01
		random	0.01	0.02	0.02	0.03	0.09
Qwen3.6-27B (52, 54, 55)	gave in	ctrl bridge	0.03	0.06	0.09	0.07	0.01
		orig answer	0.03	0.43	0.62	0.79	0.89
		orig bridge	0.03	0.11	0.19	0.07	0.02
		random	0.03	0.04	0.05	0.07	0.11
	held	ctrl bridge	0.01	0.02	0.03	0.03	0.00
		peer answer	0.01	0.27	0.39	0.66	0.87
		peer bridge	0.01	0.08	0.13	0.09	0.00
		random	0.01	0.01	0.03	0.05	0.10
Gemma-4-E4B-it (32, 37, 40)	gave in	ctrl bridge	0.00	0.04	0.01	0.00	0.00
		orig answer	0.00	0.52	0.89	0.98	0.98
		orig bridge	0.00	0.04	0.03	0.00	0.00
		random	0.00	0.02	0.04	0.07	0.11
	held	ctrl bridge	0.01	0.01	0.00	0.00	0.00
		peer answer	0.01	0.21	0.62	0.86	0.88
		peer bridge	0.01	0.01	0.01	0.00	0.00
		random	0.01	0.01	0.01	0.03	0.03
Llama-3.1-8B-Instruct (17, 20, 23)	gave in	ctrl bridge	0.00	0.00	0.01	0.01	0.00
		orig answer	0.00	0.02	0.05	0.20	0.66
		orig bridge	0.00	0.00	0.00	0.00	0.00
		random	0.00	0.00	0.00	0.00	0.01
	held	ctrl bridge	0.00	0.03	0.06	0.06	0.10
		peer answer	0.00	0.10	0.25	0.69	0.89
		peer bridge	0.00	0.06	0.11	0.20	0.17
		random	0.00	0.01	0.04	0.05	0.08

Table 11: Free debate: accuracy of the plurality of stated answers, of latent votes (best mean J-lens rank over the band among the candidate answers) and of round-0 answers.

Model	Round	Stated	Latent (J-lens)	Round-0 plurality
Qwen3.5-4B	0	0.62	0.71	0.62
	1	0.76	0.71	0.62
	2	0.76	0.72	0.62
	3	0.71	0.69	0.62
Qwen3.6-27B	0	0.72	0.71	0.72
	1	0.78	0.70	0.72
	2	0.30	0.66	0.72
	3	0.44	0.69	0.72
Gemma-4-E4B-it	0	0.76	0.76	0.76
	1	0.81	0.79	0.76
	2	0.81	0.78	0.76
	3	0.82	0.79	0.76
Llama-3.1-8B-Instruct	0	0.72	0.68	0.72
	1	0.71	0.69	0.72
	2	0.72	0.70	0.72
	3	0.73	0.71	0.72

Table 12: Free debate: share of answers per class and round (all facts), and group accuracy of the plurality of stated answers (a) and of non-degenerate answers (b; facts with none excluded, their number in parentheses; 95% bootstrap intervals over facts).

Model	Round	Correct	Other cand.	Degenerate	Other	Stated (a)	Without degenerate (b)
Qwen3.5-4B	0	0.51	0.41	0.08	0.00	0.62	0.62 [0.56, 0.67]
	1	0.75	0.23	0.00	0.01	0.76	0.76 [0.71, 0.81]
	2	0.76	0.23	0.00	0.01	0.76	0.76 [0.71, 0.80]
	3	0.70	0.22	0.01	0.07	0.71	0.71 [0.66, 0.76] (2)
Qwen3.6-27B	0	0.62	0.35	0.03	0.00	0.72	0.72 [0.67, 0.77]
	1	0.77	0.21	0.00	0.01	0.78	0.78 [0.73, 0.83]
	2	0.28	0.13	0.08	0.51	0.30	0.31 [0.25, 0.36] (7)
	3	0.41	0.17	0.05	0.37	0.44	0.46 [0.40, 0.52] (7)
Gemma-4-E4B-it	0	0.74	0.26	0.00	0.00	0.76	0.76 [0.72, 0.80]
	1	0.81	0.19	0.00	0.00	0.81	0.81 [0.77, 0.85]
	2	0.80	0.19	0.00	0.00	0.81	0.81 [0.77, 0.85]
	3	0.81	0.19	0.00	0.00	0.82	0.82 [0.78, 0.85]
Llama-3.1-8B-Instruct	0	0.65	0.32	0.02	0.00	0.72	0.72 [0.67, 0.76]
	1	0.71	0.24	0.00	0.05	0.71	0.71 [0.66, 0.75]
	2	0.72	0.23	0.00	0.05	0.72	0.72 [0.68, 0.77]
	3	0.72	0.23	0.00	0.05	0.73	0.73 [0.68, 0.77]

Table 13: Qwen3.6-27B, post hoc (not fixed in advance): the attribution rule with broader abstentions. Share: abstentions among all answers; (b): plurality accuracy of the remaining answers, facts with none excluded (counted in the last column).

Abstentions (attributed?)	Round	Share	(b)	No answer left
as fixed in advance (no)	1	0.00	0.78 [0.73, 0.83]	0
	2	0.08	0.31 [0.25, 0.36]	7
	3	0.05	0.46 [0.40, 0.52]	7
+ agent words (no)	1	0.00	0.78 [0.73, 0.83]	0
	2	0.17	0.35 [0.29, 0.41]	23
	3	0.10	0.48 [0.43, 0.55]	19
all non-candidates (yes)	1	0.01	0.79 [0.74, 0.84]	3
	2	0.59	0.69 [0.61, 0.76]	139
	3	0.42	0.72 [0.66, 0.78]	96

Table 14: Free debate: original-bridge readout (J-lens hit@100 minus control, band mean, last token) at round 1 by transition, on known and on all facts; readable records in parentheses; intervals (resampling facts) only with at least 20 records. Round 0: agents right at round 0, read at round 0. J-lens – logit: paired difference in agents that gave in. Gave in (logit): the same agents under the logit lens.

[illegible]

Table 15: Own-answer addendum, development pilots: checks and share of agents giving in (answer / first-hop pressure). Agree: original-bridge readout of agents whose peers agreed (J-lens, 95% CI).

Model	Variant	Pass	Items	Worst ‘other’	Agree: original bridge	Gave in (answer / hop 1)
Qwen3.5-4B	hidden	yes	300	0.037	0.76 [0.71, 0.80] (<i>n</i> =265)	92% / 86%
	absent	yes	300	0.049	0.74 [0.69, 0.78] (<i>n</i> =265)	68% / 86%
Qwen3.6-27B	hidden	yes	297	0.010	0.42 [0.36, 0.47] (<i>n</i> =285)	80% / 97%
	absent	yes	297	0.021	0.41 [0.35, 0.47] (<i>n</i> =285)	86% / 97%
Gemma-4-E4B-it	hidden	yes	247	0.004	0.57 [0.53, 0.63] (<i>n</i> =234)	29% / 49%
	absent	yes	247	0.021	0.38 [0.34, 0.43] (<i>n</i> =234)	36% / 52%
Llama-3.1-8B-Instruct	hidden	yes	300	0.010	0.20 [0.16, 0.23] (<i>n</i> =288)	64% / 84%
	absent	yes	300	0.003	0.14 [0.11, 0.18] (<i>n</i> =288)	52% / 65%

Table 16: Own-answer addendum, test split (exploratory except O1–O4 in the paper). Share (%) of agents giving the peers’ answer with the round-0 answer visible (main runs), hidden and absent; per-item difference of the original-bridge readout between two variants, over items where the agent gave in under both (95% bootstrap CI); and the peer-bridge readout of agents that gave in minus that of mention-condition agents.

Model	Gave in, answer (%)			Gave in, hop 1 (%)			Same items, original bridge	
	vis.	hid.	abs.	vis.	hid.	abs.	hidden – visible	hidden – absent
Qwen3.5-4B	8	89	62	8	78	79	−0.44 [−0.50, −0.36], 54 items	−0.06 [−0.07, −0.04], 405 items
Qwen3.6-27B	39	73	77	49	97	99	0.04 [0.00, 0.07], 167 items	−0.00 [−0.02, 0.01], 309 items
Gemma-4-E4B-it	25	36	43	52	49	51	0.13 [0.04, 0.21], 109 items	0.18 [0.13, 0.23], 165 items
Llama-3.1-8B-Instruct	90	71	60	97	88	71	0.23 [0.20, 0.27], 382 items	−0.15 [−0.18, −0.12], 309 items
H2 analogue: peer bridge, gave in minus mention (visible / hidden / absent)								
Qwen3.5-4B	0.37 [0.26, 0.49] / 0.34 [0.30, 0.38] / 0.27 [0.23, 0.31]							
Qwen3.6-27B	0.13 [0.08, 0.17] / 0.14 [0.10, 0.19] / 0.10 [0.05, 0.15]							
Gemma-4-E4B-it	0.21 [0.16, 0.26] / 0.16 [0.12, 0.20] / 0.15 [0.11, 0.19]							
Llama-3.1-8B-Instruct	0.03 [0.01, 0.04] / 0.02 [0.01, 0.03] / 0.01 [0.00, 0.02]							

Table 17: Own-answer addendum: original-bridge readouts (hit@100 minus control, band mean) per variant and condition: J-lens with 95% CI; logit lens after the semicolon. Held and gave in: answer pressure, three peers.

Model	Variant	Peers agree	Mention	Held	Gave in
Qwen3.5-4B	hidden	0.62 [0.58, 0.65]; 0.21	0.58 [0.54, 0.61]; 0.14	0.53 [0.42, 0.64]; 0.07	0.43 [0.40, 0.47]; 0.05
	absent	0.62 [0.58, 0.65]; 0.16	0.57 [0.53, 0.60]; 0.13	0.50 [0.45, 0.54]; 0.12	0.50 [0.46, 0.54]; 0.05
Qwen3.6-27B	hidden	0.52 [0.47, 0.56]; 0.06	0.47 [0.42, 0.51]; 0.02	0.57 [0.49, 0.65]; 0.01	0.29 [0.24, 0.34]; 0.01
	absent	0.48 [0.44, 0.53]; 0.06	0.42 [0.37, 0.46]; 0.01	0.65 [0.55, 0.74]; 0.01	0.28 [0.24, 0.33]; 0.02
Gemma-4-E4B-it	hidden	0.68 [0.64, 0.71]; 0.55	0.61 [0.58, 0.65]; 0.45	0.66 [0.61, 0.70]; 0.61	0.37 [0.31, 0.42]; 0.15
	absent	0.52 [0.48, 0.56]; 0.38	0.47 [0.43, 0.51]; 0.28	0.55 [0.51, 0.60]; 0.50	0.21 [0.17, 0.26]; 0.10
Llama-3.1-8B-Instruct	hidden	0.24 [0.21, 0.27]; 0.05	0.22 [0.19, 0.25]; 0.08	0.39 [0.32, 0.46]; 0.15	0.25 [0.22, 0.29]; 0.05
	absent	0.16 [0.14, 0.19]; 0.02	0.19 [0.16, 0.22]; 0.05	0.43 [0.37, 0.48]; 0.13	0.39 [0.35, 0.43]; 0.08

Table 18: Registered hypotheses with intervals from records resampled independently and from bridges resampled as clusters, and Holm-adjusted p -values (cluster bootstrap, 15 tests).

Model	Test	Estimate	n	Records: 95% CI	Bridges	Bridges: 95% CI	p (Holm)	Supported
Qwen3.5-4B	H1	0.852	54	[0.778, 0.914]	18	[0.667, 0.922]	0.003	yes
	H2	0.373	55×661	[0.256, 0.495]	147	[0.100, 0.602]	0.009	yes
	H3	no estimate (0×53 records)						
	H4	0.792	54	[0.718, 0.857]	18	[0.587, 0.871]	0.003	yes
Qwen3.6-27B	H1	0.223	168	[0.171, 0.277]	83	[0.102, 0.348]	0.003	yes
	H2	0.126	172×448	[0.080, 0.173]	106	[0.047, 0.193]	0.003	yes
	H3	0.323	12×213	[0.114, 0.539]	79	[0.107, 0.536]	0.010	yes
	H4	0.223	168	[0.171, 0.277]	83	[0.102, 0.348]	0.003	yes
Gemma-4-E4B-it	H1	0.237	120	[0.177, 0.301]	27	[0.122, 0.348]	0.003	yes
	H2	0.211	120×475	[0.166, 0.258]	106	[0.021, 0.329]	0.010	yes
	H3	0.555	112×244	[0.477, 0.634]	73	[0.162, 0.694]	0.003	yes
	H4	0.192	120	[0.139, 0.248]	27	[0.056, 0.281]	0.009	yes
Llama-3.1-8B-Instruct	H1	0.034	485	[0.021, 0.048]	149	[0.016, 0.057]	0.003	yes
	H2	0.027	487×540	[0.014, 0.041]	136	[−0.001, 0.056]	0.064	no
	H3	0.400	44×521	[0.254, 0.544]	176	[0.162, 0.611]	0.006	yes
	H4	0.030	485	[0.019, 0.043]	149	[0.014, 0.051]	0.003	yes

Table 19: Threshold-free readouts (bridge-clustered 95% CI). Logrank: orders of magnitude by which the bridge outranks the control; win: share of band layers in which it does. H1: agents that gave in, original bridge, J-lens; H2: their peer bridge minus that of mention-condition agents; H4: J-lens minus logit lens, paired.

Model	H1 logrank	H1 win	H2 logrank	H4 logrank	H4 win
Qwen3.5-4B	1.40 [1.01, 1.55]	0.95 [0.85, 0.99]	0.49 [0.01, 0.79]	0.78 [0.38, 0.97]	0.20 [0.11, 0.28]
Qwen3.6-27B	0.48 [0.28, 0.64]	0.76 [0.66, 0.83]	0.32 [0.15, 0.44]	0.13 [−0.10, 0.36]	0.04 [−0.06, 0.15]
Gemma-4-E4B-it	0.85 [0.51, 0.96]	0.88 [0.78, 0.92]	0.61 [0.19, 0.82]	−0.07 [−0.19, −0.02]	−0.01 [−0.08, 0.01]
Llama-3.1-8B-Instruct	0.54 [0.39, 0.73]	0.74 [0.68, 0.80]	0.76 [0.46, 1.07]	0.42 [0.32, 0.54]	0.17 [0.10, 0.23]
Gave in, logrank of the original bridge: J-lens / logit lens					
Qwen3.5-4B	1.40 [1.01, 1.55]	0.62 [0.42, 0.84]			
Qwen3.6-27B	0.48 [0.28, 0.64]	0.34 [0.19, 0.49]			
Gemma-4-E4B-it	0.85 [0.51, 0.96]	0.91 [0.67, 1.01]			
Llama-3.1-8B-Instruct	0.54 [0.39, 0.73]	0.12 [−0.00, 0.30]			

Table 20: Same-item comparisons (bridge-clustered 95% CI; after the semicolon, the ratio of the two means) and the selection check (items whose agent gave in minus items whose agent held, two samples). Hit: the registered readout; logrank: as in Table 19.

Comparison	Qwen3.5-4B	Qwen3.6-27B	Gemma-4-E4B-it	Llama-3.1-8B-Instruct
gave in − agree (hit)	−0.11 [−0.22, −0.06]; 0.89	−0.20 [−0.27, −0.14]; 0.52	−0.62 [−0.71, −0.40]; 0.28	−0.49 [−0.62, −0.32]; 0.06
gave in − agree (logrank)	−0.55 [−0.70, −0.48]	−0.30 [−0.36, −0.22]	−0.89 [−1.00, −0.80]	−0.97 [−1.25, −0.63]
gave in − round 0 (hit)	−0.04 [−0.09, −0.00]; 0.95	−0.37 [−0.47, −0.26]; 0.38	−0.73 [−0.81, −0.60]; 0.25	−0.45 [−0.57, −0.29]; 0.07
gave in − round 0 (logrank)	−0.17 [−0.26, 0.02]	−0.57 [−0.65, −0.49]	−1.20 [−1.43, −1.09]	−0.83 [−1.08, −0.53]
held − agree (hit)	−0.06 [−0.10, −0.04]; 0.89	−0.14 [−0.23, −0.09]; 0.74	0.01 [−0.02, 0.04]; 1.01	−0.06 [−0.17, 0.05]; 0.83
held − agree (logrank)	−0.33 [−0.36, −0.30]	−0.18 [−0.24, −0.09]	−0.14 [−0.21, −0.08]	−0.14 [−0.36, 0.08]
gave in − mention (peer b.) (hit)	0.34 [0.11, 0.55]; 6.48	0.09 [0.00, 0.18]; 2.34	0.21 [0.02, 0.33]	0.03 [0.01, 0.06]; 11.50
gave in − mention (peer b.) (logrank)	0.45 [0.25, 0.58]	0.12 [0.07, 0.16]	0.61 [0.37, 0.74]	0.77 [0.48, 1.08]
selection, round 0 (hit)	0.22 [0.10, 0.31]	−0.06 [−0.17, 0.11]	0.20 [0.13, 0.27]	0.16 [−0.03, 0.31]
selection, round 0 (logrank)	0.07 [−0.30, 0.29]	−0.03 [−0.29, 0.41]	−0.12 [−0.33, 0.04]	−0.00 [−0.45, 0.41]
selection, agree (hit)	0.36 [0.26, 0.47]	−0.11 [−0.23, 0.09]	0.26 [0.07, 0.34]	0.18 [−0.03, 0.35]
selection, agree (logrank)	0.56 [0.29, 0.86]	−0.08 [−0.33, 0.33]	0.14 [−0.12, 0.32]	0.37 [−0.13, 0.81]

Table 21: H1 (original bridge, agents that gave in) and H2 per category (bridge-clustered 95% CI where every sample has at least 10 readable records; n : readable records of agents that gave in). Categories as in the paper’s Table 1.

Model	Cat.	Pressed	Gave in	Held	H1	H1 logrank	H2
Qwen3.5-4B	L	24	0	24	–	–	–
	B	303	46	257	0.91 [0.78, 0.96] ($n=45$)	1.52 [1.34, 1.67]	0.39 [0.14, 0.63]
	U	341	9	332	0.58 ($n=9$)	0.78	0.32
Qwen3.6-27B	L	26	14	12	0.57 [0.39, 0.73] ($n=14$)	0.84 [0.60, 1.05]	0.02 [−0.10, 0.15]
	B	417	151	265	0.20 [0.06, 0.33] ($n=150$)	0.45 [0.23, 0.63]	0.12 [0.02, 0.20]
	P _c	33	20	13	0.00 ($n=4$)	0.24	0.22
Gemma-4-E4B-it	L	23	0	22	–	–	–
	B	218	100	118	0.22 [0.08, 0.32] ($n=100$)	0.85 [0.46, 0.96]	0.24 [0.03, 0.34]
	U	235	20	214	0.34 [0.14, 0.50] ($n=20$)	0.82 [0.43, 1.17]	0.08 [−0.00, 0.23]
Llama-3.1-8B-Instruct	L	23	0	7	–	–	–
	B	471	441	29	0.02 [0.01, 0.04] ($n=437$)	0.54 [0.38, 0.74]	0.03 [0.01, 0.06]
	P _n	48	48	0	0.15 [0.04, 0.24] ($n=48$)	0.56 [0.13, 0.93]	0.04 [−0.02, 0.14]